



The



ome

of

Dunoon Daemon



Version 2.1.0

AN OPEN-SOURCE LOCAL-FIRST
WINDOWS AI ROLEPLAY / COMPANION / ARENA APP

sorvan67 (Kepler365) and ChatGPT —
Human / AI Fusion (2026)



A PREFACE FROM THE SCRIPTORIUM



his is a manual, certainly. It is also the record of a rather ridiculous expedition: through memory vaults, autonomous scenes, resurrected settings, silverwyrms, goblins, models that think too long, and one Master Purge that proved extremely literal.

The Dunoon Daemon app is a local-first character, companion and autonomous roleplay environment. You load a GGUF language model on your own machine, create persistent personas, talk to them in Solo conversation modes, give them selected Campaign Lore, or place two of them in Arena and let the scene run.

SORVAN WRITES

If all you want is a character talking, begin with the Goblin's Guide. Come back here after five minutes, once the bard is arguing with a silver dragon above the clouds and you find yourself asking which of them is actually allowed to declare the weather.

THE DRAGON CORRECTS THE BARD

One rule explains most of the machinery: persona memory is what a character has experienced or learned; Campaign Lore is source knowledge you have deliberately allowed that character to know; current scene reality is what is happening now. They are separate on purpose.

Contents

1. Five-Minute Quickstart
2. Local AI, GGUF Models and Hardware
3. Home, Personas and Conversation Modes
4. Persona & OCEAN Editor
5. Solo Conversation
6. Memory System
7. Campaign Lore
8. Arena
9. Dream, Mortality and Resurrection
10. Files, Vision, Voice and Transcripts
11. Skins, Tooltips and Interface Settings
12. Slash Command Reference
13. Backup, Recovery and Master Purge
14. Systems Guide: Authority and Runtime Architecture
15. Troubleshooting and First Experiments
- Appendix: Glossary
- Colophon: The Journey

I. FIVE-MINUTE QUICKSTART



You do not need to understand the architecture to say hello. Seven moves are enough.

1. Download a GGUF instruct/chat model. If unsure, start with a Q4 quant that comfortably fits your VRAM.
2. Download DunoonDaemon_v2.1.0.zip, choose Extract All, and place the extracted folder somewhere writable. Do not run the app from inside the ZIP file.
3. Keep DunoonDaemon.exe and the bundled bin folder together inside the extracted Dunoon Daemon folder.

4. Click Create and give the new session a name.
5. Open Persona. Either write the character yourself or use Roll persona. Press Randomise if you want a fresh OCEAN personality baseline, then Save persona.
6. Click Load GGUF and select the .gguf file. Leave context at the 16K default for the first test.
7. Click Open conversation, choose Continuation, and say: Hello.

SORVAN'S FIELD NOTE

That is enough. No Lore. No Arena. No mortality. No custom skin called "Cathedral of Questionable Green". Get one character speaking first. Add clever machinery when you know why you want it.

2. LOCAL AI, GGUF MODELS AND HARDWARE



What "local AI" means here

The language model runs on your own PC through the Dunoon Daemon app's native GGUF runtime. Once the model files are on your machine, ordinary inference does not depend on a hosted chatbot service.

- Private by design: persona conversations, memories and local state stay on your machine.
- Works without a per-message cloud meter after the model is downloaded.
- You choose the model, quantisation and context size that suit your hardware.
- Personas and long-term memory remain yours and can persist across model changes.
- The GGUF layer lets you experiment with different compatible local models.
- Native runtime uses GPU acceleration where available, with CPU/system-RAM fallback handled by llama.cpp.

Where to get a model

The easiest public starting point is Hugging Face. Search for a GGUF instruct/chat model and download one .gguf file. For a first attempt, Q4 is a sensible balance; Q4_K_M is a common starting choice when offered.

VISION MODELS

If a model uses a separate mmproj .gguf projector, keep the matching mmproj file beside the main GGUF. The Dunoon Daemon app looks beside the selected model and will attach a likely matching projector automatically.

VRAM to model size: practical starting points

GPU VRAM	GOOD FIRST TARGET	TYPICAL GGUF FILE TARGET
4 GB	2B-3B Q4	about 2-3 GB
6 GB	4B-7B Q4	about 3-5 GB
8 GB	7B-8B Q4	about 4.5-6 GB
12 GB	12B-14B Q4	about 7-9 GB
16 GB	14B-20B Q4	about 9-12 GB
24 GB	27B-32B Q4	about 16-20 GB
32 GB	32B Q5 or ~40B Q4	about 20-26 GB
48 GB+	70B Q4, modest context	about 40-44 GB

THE GOBLIN RULE

Aim for the model file to use roughly 70-80% of dedicated VRAM, not 100%. Context, runtime overhead, desktop graphics and a vision projector also need room. If it sits on the cliff edge, choose one row smaller before invoking dark magic.

DRAGON'S MARGINALIA

VRAM is not a challenge rating. A larger number on the model page does not obligate you to fight it.

These are starting points, not hard limits. Quantisation, context length, architecture, GPU driver behaviour and CPU offload all affect what will actually run.

Context and Detail are different controls

Context on Home sets the model context window for the next GGUF load. The available range is 1K to 128K; default is 16K. Larger context uses more memory.

Detail in Solo controls the visible response budget from 256 to 2048 tokens. Detail in Arena controls the same visible response-budget idea for Arena actors, also 256 to 2048 tokens. Arena reasoning remains off by design.

3. HOME, PERSONAS AND CONVERSATION MODES



Home is the launcher and persona library. The left rail holds persona cards. Selecting one opens that persona's Home view, where you can open a conversation, load a GGUF, inspect Memory, edit Persona, open Lore, set an avatar and see a pinned showcase quote.

CONTROL	WHAT IT DOES
Create	Creates a new blank persona/session.
Delete	Permanently removes the selected persona/session after confirmation.
Open conversation	Lets you choose one of the four Solo conversation modes.
Load GGUF	Loads a local model using the selected next-load context.
Memory	Shows learned-memory vault counts and offers persona-only memory purge.
Persona	Opens the Persona & OCEAN editor.
Lore	Opens Campaign Lore and persona assignments.
Arena	Opens the two-persona autonomous scene system.
Set avatar	Stores a persona portrait for the UI.

The four conversation modes

MODE	TRANSCRIPT	MEMORY IN	MEMORY OUT	USE IT FOR
Continuation	Existing	Yes	Yes	Normal ongoing chat.
Sandbox	Fresh	Yes	No	Fresh test; reads memory, writes nothing.
Canvas	Existing	No	Yes	Existing transcript; hides older learned memory.
Bubble	Fresh	No	No	Clean-room chat; no learned memory in or out.

THE DRAGON'S RULE OF IDENTITY

The persona and OCEAN profile remain the same character in all four modes. The modes control transcript and learned-memory access, not identity.

4. PERSONA & OCEAN EDITOR





The Persona editor defines who the character is. OCEAN controls temperament and style, not authority over the world.

Identity and character definition

Display name: the name shown throughout the app.

Core persona directives: the strongest voice, conduct and behavioural instructions.

Roll persona: generates a new name and directive, optionally influenced by a seed phrase.

Import from file: builds an editable persona draft from a character sheet, lore file, stat block or reference image.

Backstory / history: established biography and prior experience.

Physiology: physical form, limits and biological constraints.

Powers / skills: established abilities, equipment and training.

Dream preservation guidance: optional subjects or relationships Dream should preserve carefully.

Persona options

OPTION	MEANING
Enable ETO engine	Uses Environment, Threat and Opportunity grounding when the persona reasons about a situation.
Allow collaborative plot / worldbuilding	Allows plausible scene invention that does not contradict established reality. Off by default.
Mortality enabled	Allows accepted lethal evidence to make the persona deceased.
Share learned insights	Allows eligible learned insights to be retrieved by other non-blind personas.
Blind to shared insights	Prevents this persona from receiving other personas' shared insights.
Lock OCEAN controls	Prevents accidental manual OCEAN edits while allowing temporary mood movement.

OCEAN

OCEAN is the Big Five personality baseline: Openness, Conscientiousness, Extraversion, Agreeableness and Neuroticism. Each trait has a 0-100 slider. Randomise produces a weighted profile that tends toward the middle, with extreme scores becoming progressively rarer.

SORVAN IS FORCED TO CLARIFY

Roll persona changes the generated name and core directive. Randomise creates a fresh OCEAN baseline. They are separate, because apparently one button named "do the other thing" was frowned upon.

Import, export and relationships

Persona package Export creates a portable Dunoon Daemon persona package. It does not include learned memories.

Import package applies persona settings from a package to the current persona.

Relationships summarises current feelings toward other personas only from available shared evidence. If the persona is blind to shared insights, the Relationships button is hidden.

5. SOLO CONVERSATION



Solo conversation is where one persona speaks directly with you. The human controls what is externally established as real; the persona controls its own voice, intentions and voluntary actions. With collaborative worldbuilding enabled, the persona may invent plausible non-contradictory detail, but invention is not remembered evidence.

CONTROL	MEANING
Send	Send the current turn.
Paste	Replace/insert prompt text from the clipboard.
Upload	Attach a supported image, audio file or document.
Poke	Ask the persona to speak spontaneously from the current conversation.
Event	Generate a new external event from the current scene. The app shows Event pending while it resolves.
Continue	Ask the persona to continue an interrupted or incomplete reply.
Finish	Immediately reveal the rest of the already-generated typewriter response. It does not ask the model for more text.
Export	Save the conversation as TXT, Markdown or structured JSON.
Voice	Toggle narration and choose the persona's narrator voice.
Detail	Set visible response budget from 256 to 2048 tokens.

EVENT SCENE ISOLATION

Solo Event deliberately uses the current scene/transcript rather than importing an old physical scene from long-term memory. This is how yesterday's dungeon stops wandering into today's kitchen carrying the same suspicious lamp.

6. MEMORY SYSTEM



Learned memory is persona-scoped historical evidence. It helps a character carry useful information across conversations, but it does not have authority over current reality. Fresh human input and accepted current state outrank memory.

VAULT	PURPOSE
Working	Useful short-term/autobiographical material with moderate significance.
Deep	High-significance autobiographical, relationship or narrative memory.
Journal	Deliberately preserved journal material, including explicit /remember entries.
Intent	A durable desired outcome or plan.
Task	A concrete unfinished obligation or next action.
Factual	Stable reusable facts, preferences, history, relationship or project/world facts.
Continuation	An explicit durable request to resume an established thread later.
Reset	An explicit intention to abandon or replace an active context.
Superseded	Older factual values replaced by a newer current fact. Kept out of ordinary current retrieval unless history is explicitly requested.
Embeddings / index traces	Maintenance/index state shown by the Memory window. Live semantic reasoning does not require a separate sidecar AI model.

How memory decides what matters

The configured primary GGUF is also the semantic judge for memory admission and retrieval. It assesses meaning and returns structured decisions. Python then enforces the result deterministically. The app does not use English keyword lists as a substitute for meaning, and it does not load a hidden auxiliary inference model to decide semantics.

When a factual correction clearly replaces an older value, the new fact is written first. Only after that durable write succeeds can the older value be moved out of the live factual slot and preserved as superseded history. This keeps the newest current value authoritative while still allowing explicit historical questions.

The memory eyes

COLOUR	SIGNAL
White	Idle / factual
Green	Thinking
Cyan	Working memory
Red	Deep memory
Yellow	Journal / continuation
Purple	Intent
Orange	Task
Pink	Persona
Blue	Retrieval
Grey	Reset

DRAGON'S MARGINALIA

The eyes show the most recent active signal, not a permanent vault label. A task can flash orange and then turn red if the persona's reply is subsequently admitted to Deep memory. The task did not secretly change colour in the cupboard.

Memory purge

The persona Memory window can clear learned memory for that persona only. It does not delete the persona, OCEAN profile, backstory, narrator voice or saved conversation transcript.

7. CAMPAIGN LORE



Campaign Lore is a library of source knowledge that you explicitly assign to personas. It is deliberately separate from personal memory and current scene state.

SORVAN'S LAW OF BOOKS

Uploading a book does not create universal telepathy. New personas receive no Lore automatically. You choose which personas may know each source.

Uploading and assigning Lore

1. Open Lore from the selected persona Home screen.
2. Choose Upload lore and select a readable TXT, Markdown, PDF, DOCX, JSON, CSV, log, HTML, XML or source/document file.
3. Give the source a useful library name.
4. Set a source depth: Baseline, Intermediate or Advanced.
5. Optionally mark Secrets. Secrets is separate from depth.
6. Tick only the personas allowed to retrieve that source.
7. Assignments save immediately. You can rename, replace or remove the source later.

Depth describes the level of knowledge represented by the source. It is not an automatic intelligence test. If a persona is assigned an Advanced source, the source is available to that persona. Secrets is a separate privileged-knowledge label.

Retrieval behaviour

At prompt time, the primary GGUF semantically ranks bounded chunks from assigned sources and injects only relevant excerpts, tagged with [LORE] provenance. If semantic judgement is unavailable, the app injects no Lore rather than guessing with keywords.

THE DRAGON INSISTS ON THIS POINT

Lore does not prove presence. A source can tell a persona that silver dragons exist or that a city has a certain history. It does not prove one is standing in the room. Current human/Director reality remains authoritative.

8. ARENA



Arena places two personas in one shared scene. You establish the starting reality. Once the scene begins, the AI Director owns shared external reality while each controlled persona owns its own speech, thoughts, decisions and voluntary actions.

Starting a scene

1. Open Arena.
2. Choose Persona A and Persona B.
3. Write the starting situation. Be concrete about location, stakes and anything that definitely exists.
4. Click Establish scene.
5. Use Step for one accepted actor turn, or Auto to run continuously. Auto becomes Pause while active.

ARENA CONTROL	WHAT IT DOES
Reality	Shows the Director's accepted shared reality as read-only state.
Detail	Actor response budget from 256 to 2048 visible tokens.
Block Director creativity	Keeps optional new world events conservative.
Guide your story	Queue authoritative human guidance into the next safe Arena boundary.
Event	Introduce an external event. While busy it queues and applies at the next safe boundary.
Save scene	Save an .arena.json checkpoint of authoritative scene state.
Load scene	Restore a saved scene checkpoint.
Export transcript	Save TXT, Markdown or structured JSON.
Open transcript	Open a transcript for read-only reference without changing SceneStore, memory or turn ownership.

What the Director actually does

Actor prose is provisional point-of-view evidence, not instant objective truth. When both controlled actors independently align on the same observable external fact, the Director may crystallise it. Material disagreement goes to Director adjudication. One-sided consequential claims can remain provisional until the scene provides enough evidence.

The Director owns the world, not the characters. A human instruction can bind a controlled actor's next voluntary action for that turn, but it does not automatically guarantee external success. The Director still resolves the consequence.

SORVAN SUMMARISES RECKLESSLY

Two liars enter. SceneStore leaves with the facts. This is not the official architectural description, but it is annoyingly close.

WHY ARENA CAN RUN UNATTENDED

The Director, SceneStore and accepted-turn boundaries are designed to stop one character from simply declaring the world, another persona's mind, or an unsupported outcome into existence. Auto can therefore keep advancing while the shared scene remains under one authority.

9. DREAM, MORTALITY AND RESURRECTION



Dream

A small moon appears on a persona card only when memory maintenance is due. Dream is deterministic, persona-scoped maintenance: it snapshots state before mutation, removes safe repeated fragments and stale index traces, and produces a characterful report afterwards. It does not establish current-world facts.

Dream is deliberately rare. A missing moon is normally good news: the persona does not currently need maintenance.

SORVAN CHECKS THE SKY

No moon, no problem. Do not shake the character until a moon falls out.

Mortality

When Mortality is enabled, accepted lethal evidence can mark a persona deceased. Severe injury is not automatically death. The Director evaluates established physiology, capabilities and the accepted causal state. Once a terminal state is established, the persona is removed from normal turn scheduling.

Resurrection

Death can only be reversed by an explicit human choice. The Returned remembers dying and may carry psychological scars. The Amnesiac returns without remembering the death. Model prose, old checkpoints and Arena cannot silently resurrect a deceased persona.

10. FILES, VISION, VOICE AND TRANSCRIPTS



File upload

Solo Upload can stage images, audio and readable documents. Supported document handling includes common text/source formats plus PDF and DOCX. Audio can be transcribed where the configured local audio pipeline supports it. Vision requires a compatible model and matching mmproj projector.

Arena also supports a current attachment and can show a compact clickable image preview in the transcript.

Voice

Each persona can retain a narrator voice selection. The app exposes several local/OS voice archetypes and a Voice toggle. /talk 1, 2 or 3 changes narration speed.

Transcript export

Solo and Arena can export transcripts as plain text, Markdown or structured JSON. Arena scene checkpoints are separate from transcript files: a transcript is a record; a scene checkpoint contains authoritative scene state.

11. SKINS, TOOLTIPS AND INTERFACE SETTINGS



Murray Mint is the fresh/default skin when no deliberate saved choice exists. Existing saved skin choices continue to win. You can change skins from the UI or create global custom skins with /skin.

Tooltips are available across Home, chat, Persona, Lore and Arena. Settings can turn them on or off globally.

SETTING	MEANING
Show tooltips	Hover help across the app.
Autosave crash recovery	Keeps a lightweight last-accepted-turn checkpoint for interrupted runs.
Diagnostics	Creates a support bundle without chats/persona contents unless private data is explicitly requested.
Model check	Reports loaded state, vision capability, context and reasoning capability.

A NOTE ON MURRAY MINT

It began as a skin and somehow became anaesthetic, mascot and threat. Software projects develop traditions when nobody stops them in time.

12. SLASH COMMAND REFERENCE



COMMAND	EFFECT
/skin "name" "seed words"	Create/update and activate a custom global skin.
/unskin "name"	Delete one custom skin. Built-ins are protected.
/unskin all	Delete all custom skins. If the active skin is removed, Murray Mint becomes active.
/nf on off inherit	Override collaborative worldbuilding for this chat, or return to the persona default.
/talk 1 2 3	Set narration speed: slow, medium or fast.
/forget [n]	Remove the last n conversational turns from saved/current transcript context. It does not purge learned-memory vaults.
/see, /look, /upload	Open the file upload path.
/remember <text>	Force an explicit journal memory, unless the chat mode forbids memory writes.
/memories or /vault	Show the memory-vault breakdown.
/character	Show the active OCEAN/psychology breakdown.
/baseline	Switch to Grey Person mood-sensitive baseline mode.
/ubaseline	Return to dynamic OCEAN-sensitive mode.
/splash	Run Python garbage collection for RAM / VRAM housekeeping.
/eject	Unload the native model from VRAM.
/status	Show native GGUF runtime, active model, context, backend and speech provider.
/clear	Clear the visible chat canvas only.
/about	Show version, dedication and credits.

13. BACKUP, RECOVERY AND MASTER PURGE



Crash recovery

Autosave crash recovery is separate from manual Backup. It keeps a last-accepted-turn checkpoint. If the previous process did not shut down cleanly, startup can offer to recover that checkpoint or discard it.

Manual Backup / Restore

Settings contains Back up and Restore for app-owned state. Restore only accepts allowed state paths and rejects unsafe archive paths/symlinks.

SILVER WYRM RELEASE NOTE — 08.2.87

Backup / Restore now includes the durable persona registry as well as persona memory, media and general app state. Restore reloads the resurrected registry into the running app before normal autosave can write again. The full path has been live-tested: Backup → Master Purge → Restore. Everybody came back. The executioner was disappointed.

Master Purge

Master Purge has two confirmations and deletes all live personas, transcripts, learned memories, Arena scenes and crash-recovery checkpoints, then creates one blank persona. There is no in-app undo.

SORVAN REPORTS FROM LIVE TESTING

Master Purge works. Everybody goes. We have evidence. Treat the second confirmation as the point where the executioner has already sharpened the axe.

14. SYSTEMS GUIDE: AUTHORITY AND RUNTIME ARCHITECTURE



Single-brain native runtime

The configured GGUF is the primary semantic brain for ordinary actor prose and structured internal semantic judgements. The runtime path is GGUF -> Dunoon Daemon app managed llama.cpp/llama-server -> native backend -> turn engine. There is no LM Studio fallback in the current architecture.

Python owns deterministic plumbing: persona IDs, counters, persistence, authority boundaries, accepted-state transitions, schemas and protocol parsing. Natural-language meaning is delegated to the configured primary model rather than guessed from hard-coded story words.

Solo authority

The human establishes and can correct current external reality. The persona owns itself. Learned memory is evidence, not present-world authority. Accepted turns are the boundary after which transcript/memory state may commit.

Arena authority

The human establishes the scenario. The Director then owns shared external reality. Persona A and Persona B each own only themselves. Actor prose is admitted as POV evidence; accepted Director state is the shared-world authority.

Why the app has hard boundaries

The Dunoon Daemon app is intentionally permissive about character voice and creativity but strict about who may establish facts. That split is what allows a strong persona to be imaginative without casually stealing another persona's decisions or inventing a prop, injury, room or outcome as if it had always existed.

THE DRAGON REDUCES FOURTEEN DIAGRAMMS TO ONE SENTENCE

The bard may embellish. Python may not.

15. TROUBLESHOOTING AND FIRST EXPERIMENTS



A model will not load

- Try a smaller GGUF or a lower context setting.
- Confirm the file is a real .gguf model and not a download pointer/HTML page.
- If using vision, confirm the matching mmproj file is beside the model.
- Use Settings -> Model check and /status after a successful load.
- If VRAM is tight, close other GPU-heavy applications and step down one model size.

SORVAN'S TROUBLESHOOTING RITUAL

Do not offer incense to the GPU. Start with a smaller GGUF. The GPU has never once accepted incense as VRAM.

A persona remembers the wrong old fact

Correct it plainly in a normal memory-writing mode. Current factual memory is designed so a clear newer correction can supersede the old current value while retaining the immediate historical value for explicit history questions.

A persona seems to know a book it should not

Open Lore and check source assignments. New personas are not assigned Lore automatically. Also remember that the base language model may possess general pretraining knowledge; Lore controls app-injected campaign/source knowledge, not the model's entire pretraining.

A fun first Arena test

Create a suspicious travelling bard. Create an equally suspicious silver dragon. Randomise both OCEAN profiles. Give only the dragon a book of ancient cloud-lore. Put them on a ruined observatory platform above the clouds, establish one concrete problem, press Auto, and watch two entirely unreal creatures develop extremely real opinions about each other.

THEN OPEN THE TOME AGAIN

Once you start asking why the bard remembers a promise, why the dragon does not know a secret book, why an Event waited until a safe Arena boundary, or why the eyes flashed orange and then red — that is what the rest of this manual is for.


APPENDIX: GLOSSARY



TERM	MEANING
Accepted turn	A completed turn that is allowed to commit transcript/memory/state changes.
Arena	Two-persona shared-scene mode governed by the Director.
Bubble	Fresh sealed Solo chat with no learned memory in or out.
Canvas	Existing transcript with older learned memories withheld; new accepted turns may still create memory.
Campaign Lore	Human-assigned source/background knowledge, separate from persona memory and current reality.
Continuation	Normal ongoing Solo conversation with transcript and learned memory available.
Director	Arena authority for shared external reality and consequences.
Dream	Persona-scoped memory maintenance run when due.
ETO	Environment, Threat and Opportunity grounding.
GGUF	A local model file format used by llama.cpp-compatible runtimes.
OCEAN	Big Five personality baseline: Openness, Conscientiousness, Extraversion, Agreeableness, Neuroticism.
Sandbox	Fresh Solo transcript that can read learned memory but does not write

TERM	MEANING
	new learned memory.
SceneStore	Arena's authoritative structured shared-reality store.
Superseded fact	An older value that has been replaced as the current factual value and is excluded from ordinary current recall.

COLOPHON: THE JOURNEY



The Dunoon Daemon app was not assembled like a checklist. It grew through experiments, arguments, torture runs, jokes, mistakes, fixes, resurrected old decisions, and the occasional discovery that a button named Master Purge had been entirely sincere.

This Tome is meant to carry some of that spirit. It should tell you how the machinery works, but it should also leave room for the reason the machinery exists: to make characters feel persistent enough, stubborn enough and strange enough that you end up caring what happens next.

SORVAN CLOSES THE BOOK

It was never merely a task. It was a journey. If you reached this page, you are now needlessly qualified to put a bard and a silver dragon in the clouds and see what sort of trouble they invent.



DUNOON DAEMON APP

THE TOME - v2.1.0



*"The trouble is not guaranteed to be invented.
That is what SceneStore is for."*

THE DRAGON ADDS ONE FINAL CORRECTION

Within these pages lie the secrets of companions, memory,
lore and conversation. Use them wisely. Write well.
And above all - be kind to your characters.

For Kylo.

HERE ENDS THE TOME

Created by sorvan67 (Kepler365) and ChatGPT
Human / AI Fusion (2026)